

LanguageManual DML

Created by Confluence Administrator, last modified by Lefty Leverenz on Aug 28, 2015

Hive Data Manipulation Language

- [Hive Data Manipulation Language](#)
 - [Loading files into tables](#)
 - [Syntax](#)
 - [Synopsis](#)
 - [Notes](#)
 - [Inserting data into Hive Tables from queries](#)
 - [Syntax](#)
 - [Synopsis](#)
 - [Notes](#)
 - [Dynamic Partition Inserts](#)
 - [Example](#)
 - [Additional Documentation](#)
 - [Writing data into the filesystem from queries](#)
 - [Syntax](#)
 - [Synopsis](#)
 - [Notes](#)
 - [Inserting values into tables from SQL](#)
 - [Syntax](#)
 - [Synopsis](#)
 - [Examples](#)
 - [Update](#)
 - [Syntax](#)
 - [Synopsis](#)
 - [Notes](#)
 - [Delete](#)
 - [Syntax](#)
 - [Synopsis](#)
 - [Notes](#)

There are multiple ways to modify data in Hive:

- [LOAD](#)
- [INSERT](#)
 - [into Hive tables from queries](#)
 - [into directories from queries](#)
 - [into Hive tables from SQL](#)
- [UPDATE](#)
- [DELETE](#)

[EXPORT](#) and [IMPORT](#) commands are also available (as of [Hive 0.8](#)).

Loading files into tables

Hive does not do any transformation while loading data into tables. Load operations are currently pure copy/move operations that move datafiles into locations corresponding to Hive tables.

Syntax

```
LOAD DATA [LOCAL] INPATH 'filepath' [OVERWRITE] INTO TABLE tablename [PARTITION (part1=part1val,
```

Synopsis

Load operations are currently pure copy/move operations that move datafiles into locations corresponding to Hive tables.

- *filepath* can be:
 - a relative path, such as `project/data1`
 - an absolute path, such as `/user/hive/project/data1`
 - a full URI with scheme and (optionally) an authority, such as `hdfs://namenode:9000/user/hive/project/data1`
- The target being loaded to can be a table or a partition. If the table is partitioned, then one must specify a specific partition of the table by specifying values for all of the partitioning columns.
- *filepath* can refer to a file (in which case Hive will move the file into the table) or it can be a directory (in which case Hive will move all the files within that directory into the table). In either case, *filepath* addresses a set of files.
- If the keyword LOCAL is specified, then:
 - the load command will look for *filepath* in the local file system. If a relative path is specified, it will be interpreted relative to the user's current working directory. The user can specify a full URI for local files as well - for example: `file:///user/hive/project/data1`
 - the load command will try to copy all the files addressed by *filepath* to the target filesystem. The target file system is inferred by looking at the location attribute of the table. The copied data files will then be moved to the table.
- If the keyword LOCAL is *not* specified, then Hive will either use the full URI of *filepath*, if one is specified, or will apply the following rules:
 - If scheme or authority are not specified, Hive will use the scheme and authority from the hadoop configuration variable `fs.default.name` that specifies the Namenode URI.
 - If the path is not absolute, then Hive will interpret it relative to `/user/<username>`
 - Hive will *move* the files addressed by *filepath* into the table (or partition)
- If the OVERWRITE keyword is used then the contents of the target table (or partition) will be deleted and replaced by the files referred to by *filepath*; otherwise the files referred by *filepath* will be added to the table.
 - Note that if the target table (or partition) already has a file whose name collides with any of the filenames contained in *filepath*, then the existing file will be replaced with the new file.

Notes

- *filepath* cannot contain subdirectories.
- If the keyword LOCAL is not given, *filepath* must refer to files within the same filesystem as the table's (or partition's) location.
- Hive does some minimal checks to make sure that the files being loaded match the target table. Currently it checks that if the table is stored in sequencefile format, the files being loaded are also sequencefiles, and vice versa.
- A bug that prevented loading a file when its name includes the "+" character is fixed in release 0.13.0 ([HIVE-6048](#)).
- Please read [CompressedStorage](#) if your datafile is compressed.

Inserting data into Hive Tables from queries

Query Results can be inserted into tables by using the insert clause.

Syntax

Standard syntax:

```
INSERT OVERWRITE TABLE tablename1 [PARTITION (partcol1=val1, partcol2=val2 ...) [IF NOT EXISTS]] select_statement1
INSERT INTO TABLE tablename1 [PARTITION (partcol1=val1, partcol2=val2 ...)] select_statement1
```

Hive extension (multiple inserts):

```
FROM from_statement
INSERT OVERWRITE TABLE tablename1 [PARTITION (partcol1=val1, partcol2=val2 ...) [IF NOT EXISTS]] select_statement1
[INSERT OVERWRITE TABLE tablename2 [PARTITION ... [IF NOT EXISTS]] select_statement2]
```

```
[INSERT INTO TABLE tablename2 [PARTITION ...] select_statement2] ...;
FROM from_statement
INSERT INTO TABLE tablename1 [PARTITION (partcol1=val1, partcol2=val2 ...)] select_st
[INSERT INTO TABLE tablename2 [PARTITION ...] select_statement2]
[INSERT OVERWRITE TABLE tablename2 [PARTITION ... [IF NOT EXISTS]] select_statement2]

Hive extension (dynamic partition inserts):
INSERT OVERWRITE TABLE tablename PARTITION (partcol1[=val1], partcol2[=val2] ...) sel
INSERT INTO TABLE tablename PARTITION (partcol1[=val1], partcol2[=val2] ...) select_s
```

Synopsis

- INSERT OVERWRITE will overwrite any existing data in the table or partition
 - unless IF NOT EXISTS is provided for a partition (as of Hive [0.9.0](#)).
- INSERT INTO will append to the table or partition, keeping the existing data intact. (Note: INSERT INTO syntax is only available starting in version 0.8.)
 - As of Hive [0.13.0](#), a table can be made **immutable** by creating it with [TBLPROPERTIES](#) ("immutable"="true"). The default is "immutable"="false".
 INSERT INTO behavior into an immutable table is disallowed if any data is already present, although INSERT INTO still works if the immutable table is empty. The behavior of INSERT OVERWRITE is not affected by the "immutable" table property.
 An immutable table is protected against accidental updates due to a script loading data into it being run multiple times by mistake. The first insert into an immutable table succeeds and successive inserts fail, resulting in only one set of data in the table, instead of silently succeeding with multiple copies of the data in the table.
- Inserts can be done to a table or a partition. If the table is partitioned, then one must specify a specific partition of the table by specifying values for all of the partitioning columns. If [hive.typecheck.on.insert](#) is set to true, these values are validated, converted and normalized to conform to their column types (Hive [0.12.0](#) onward).
- Multiple insert clauses (also known as *Multi Table Insert*) can be specified in the same query.
- The output of each of the select statements is written to the chosen table (or partition). Currently the OVERWRITE keyword is mandatory and implies that the contents of the chosen table or partition are replaced with the output of corresponding select statement.
- The output format and serialization class is determined by the table's metadata (as specified via DDL commands on the table).
- As of [Hive 0.14](#), if a table has an OutputFormat that implements AcidOutputFormat and the system is configured to use a [transaction](#) manager that implements ACID, then INSERT OVERWRITE will be disabled for that table. This is to avoid users unintentionally overwriting transaction history. The same functionality can be achieved by using [TRUNCATE TABLE](#) (for non-partitioned tables) or [DROP PARTITION](#) followed by INSERT INTO.
- As of Hive [1.1.0](#) the TABLE keyword is optional.
- As of Hive [1.2.0](#) each INSERT INTO T can take a column list like INSERT INTO T (z, x, c1). See Description of [HIVE-9481](#) for examples.

Notes

- Multi Table Inserts minimize the number of data scans required. Hive can insert data into multiple tables by scanning the input data just once (and applying different query operators) to the input data.
- Starting with [Hive 0.13.0](#), the select statement can include one or more common table expressions (CTEs) as shown in the [SELECT syntax](#). For an example, see [Common Table Expression](#).

Dynamic Partition Inserts

Version information

This information reflects the situation in Hive 0.12; dynamic partition inserts were added in Hive 0.6.

In the dynamic partition inserts, users can give partial partition specifications, which means just specifying the list of partition column names in the `PARTITION` clause. The column values are optional. If a partition column value is given, we call this a static partition, otherwise it is a dynamic partition. Each dynamic partition column has a corresponding input column from the select statement. This means that the dynamic partition creation is determined by the value of the input column. The dynamic partition columns must be **specified last** among the columns in the `SELECT` statement and **in the same order** in which they appear in the `PARTITION()` clause.

Dynamic Partition inserts are disabled by default. These are the relevant configuration properties for dynamic partition inserts:

Configuration property	Default	Note
<code>hive.exec.dynamic.partition</code>	false	Needs to be set to true to enable dynamic partition inserts
<code>hive.exec.dynamic.partition.mode</code>	strict	In strict mode, the user must specify at least one static partition in case the user accidentally overwrites all partitions, in nonstrict mode all partitions are allowed to be dynamic
<code>hive.exec.max.dynamic.partitions.pernode</code>	100	Maximum number of dynamic partitions allowed to be created in each mapper/reducer node
<code>hive.exec.max.dynamic.partitions</code>	1000	Maximum number of dynamic partitions allowed to be created in total
<code>hive.exec.max.created.files</code>	100000	Maximum number of HDFS files created by all mappers/reducers in a MapReduce job
<code>hive.error.on.empty.partition</code>	false	Whether to throw an exception if dynamic partition insert generates empty results

Example

```
FROM page_view_stg pvs
INSERT OVERWRITE TABLE page_view PARTITION(dt='2008-06-08', country)
SELECT pvs.viewTime, pvs.userid, pvs.page_url, pvs.referrer_url, null, null, p
```

Here the country partition will be dynamically created by the last column from the `SELECT` clause (i.e. `pvs.cnt`). Note that the name is not used. In nonstrict mode the dt partition could also be dynamically created.

Additional Documentation

- [Design Document](#)
 - [Original design doc](#)
 - [HIVE-936](#)
- [Tutorial: Dynamic-Partition Insert](#)
- [HCatalog Dynamic Partitioning](#)
 - [Usage with Pig](#)
 - [Usage from MapReduce](#)

Writing data into the filesystem from queries

Query results can be inserted into filesystem directories by using a slight variation of the syntax above:

Syntax

Standard syntax:

```
INSERT OVERWRITE [LOCAL] DIRECTORY directory1
  [ROW FORMAT row_format] [STORED AS file_format] (Note: Only available starting with
  SELECT ... FROM ...
```

Hive extension (multiple inserts):

```
FROM from_statement
INSERT OVERWRITE [LOCAL] DIRECTORY directory1 select_statement1
[INSERT OVERWRITE [LOCAL] DIRECTORY directory2 select_statement2] ...
```

row_format

```
: DELIMITED [FIELDS TERMINATED BY char [ESCAPED BY char]] [COLLECTION ITEMS TERMINATED BY char]
  [MAP KEYS TERMINATED BY char] [LINES TERMINATED BY char]
  [NULL DEFINED AS char] (Note: Only available starting with Hive 0.13)
```

Synopsis


- Directory can be a full URI. If scheme or authority are not specified, Hive will use the scheme and authority from the hadoop configuration variable `fs.default.name` that specifies the Namenode URI.
- If LOCAL keyword is used, Hive will write data to the directory on the local file system.
- Data written to the filesystem is serialized as text with columns separated by ^A and rows separated by newlines. If any of the columns are not of primitive type, then those columns are serialized to JSON format.

Notes

- INSERT OVERWRITE statements to directories, local directories, and tables (or partitions) can all be used together within the same query.
- INSERT OVERWRITE statements to HDFS filesystem directories are the best way to extract large amounts of data from Hive. Hive can write to HDFS directories in parallel from within a map-reduce job.
- The directory is, as you would expect, OVERWRITTEN; in other words, if the specified path exists, it is clobbered and replaced with the output.
- As of Hive 0.11.0 the separator used can be specified; in earlier versions it was always the ^A character (\001). However, custom separators are only supported for LOCAL writes in Hive versions 0.11.0 to 1.1.0 – this bug is fixed in version 1.2.0 (see [HIVE-5672](#)).
- In [Hive 0.14](#), inserts into ACID compliant tables will deactivate vectorization for the duration of the select and insert. This will be done automatically. ACID tables that have data inserted into them can still be queried using vectorization.

Inserting values into tables from SQL

The INSERT...VALUES statement can be used to insert data into tables directly from SQL.


Version Information
 INSERT...VALUES is available starting in [Hive 0.14](#).

 Inserting values from SQL statements can only be performed on tables that support ACID. See [Hive Transactions](#) for details.

Syntax

Standard Syntax:

```
INSERT INTO TABLE tablename [PARTITION (partcol1[=val1], partcol2[=val2] ...)] VALUES
```

Where values_row is:

```
( value [, value ...] )
```

where a value is either **null** or any valid SQL literal

Synopsis

- Each row listed in the VALUES clause is inserted into table *tablename*.
- Values must be provided for every column in the table. The standard SQL syntax that allows the user to insert values into only some columns is not yet supported. To mimic the standard SQL, nulls can be provided for columns the user does not wish to assign a value to.
- Dynamic partitioning is supported in the same way as for [INSERT...SELECT](#).
- If the table being inserted into supports [ACID](#) and a transaction manager that supports ACID is in use, this operation will be auto-committed upon successful completion.
- Insert, update, delete operations are not supported on tables that are sorted (tables created with the [SORTED BY](#) clause).
- Hive does not support literals for complex types (array, map, struct, union), so it is not possible to use them in INSERT INTO...VALUES clauses. This means that the user cannot insert data into a complex datatype column using the INSERT INTO...VALUES clause.

Examples

```
CREATE TABLE students (name VARCHAR(64), age INT, gpa DECIMAL(3, 2))
  CLUSTERED BY (age) INTO 2 BUCKETS STORED AS ORC;
```

```
INSERT INTO TABLE students
  VALUES ('fred flintstone', 35, 1.28), ('barney rubble', 32, 2.32);
```

```
CREATE TABLE pageviews (userid VARCHAR(64), link STRING, came_from STRING)
  PARTITIONED BY (datestamp STRING) CLUSTERED BY (userid) INTO 256 BUCKETS STORED AS ORC;
```

```
INSERT INTO TABLE pageviews PARTITION (datestamp = '2014-09-23')
  VALUES ('jsmith', 'mail.com', 'sports.com'), ('jdoe', 'mail.com', null);
```

```
INSERT INTO TABLE pageviews PARTITION (datestamp)
  VALUES ('tjohnson', 'sports.com', 'finance.com', '2014-09-23'), ('tlee', 'finance.com', null, null);
```

Update



Version Information

UPDATE is available starting in [Hive 0.14](#).

Updates can only be performed on tables that support ACID. See [Hive Transactions](#) for details.

Syntax

Standard Syntax:

```
UPDATE tablename SET column = value [, column = value ...] [WHERE expression]
```

Synopsis

- The referenced column must be a column of the table being updated.
- The value assigned must be an expression that Hive supports in the select clause. Thus arithmetic operators, UDFs, casts, literals, etc. are supported. Subqueries are not supported.
- Only rows that match the WHERE clause will be updated.
- Partitioning columns cannot be updated.
- Bucketing columns cannot be updated.
- In Hive 0.14, upon successful completion of this operation the changes will be auto-committed.

Notes

- Vectorization will be turned off for update operations. This is automatic and requires no action on the part of the user. Non-update operations are not affected. Updated tables can still be queried using vectorization.
- In version 0.14 it is recommended that you set [hive.optimize.sort.dynamic.partition=false](#) when doing updates, as this produces more efficient execution plans.

Delete**Version Information**

DELETE is available starting in [Hive 0.14](#).

Deletes can only be performed on tables that support ACID. See [Hive Transactions](#) for details.

Syntax**Standard Syntax:**

```
DELETE FROM tablename [WHERE expression]
```

Synopsis

- Only rows that match the WHERE clause will be deleted.
- In Hive 0.14, upon successful completion of this operation the changes will be auto-committed.

Notes

- Vectorization will be turned off for delete operations. This is automatic and requires no action on the part of the user. Non-delete operations are not affected. Tables with deleted data can still be queried using vectorization.
- In version 0.14 it is recommended that you set [hive.optimize.sort.dynamic.partition=false](#) when doing deletes, as this produces more efficient execution plans.

[Muhammad Ahsan](#) likes this

No labels

Confluence today.

This Confluence installation runs a Free Gliffy License - Evaluate
the Gliffy Confluence Plugin for your Wiki!