

SPA Conference 2016

Real-World Big Data in Action

Nick Rozanski
Eoin Woods
Chris Cooper-Bland



Number 2 in an occasional series

OS X and Linux User Setup

Create a spa16 User in OS X

- Use System Preferences -> Users and Groups
- Set the Account Name to `spa16`
- Make the user an Administrator
- Log out of OS X, and log back in again as the `spa16` user

Create a spa16 User in Linux

- Enter the following from a Linux command prompt:
`sudo adduser spa16`
- Make the user an Administrator
`gpasswd -a spa16 sudo`
- Log back in as the `spa16` user

Set up Passphraseless SSH (OS X, Linux)

Spark and SSH

- SSH is Secure Shell, a cryptographically secure way of running services over an insecure network (for example, logging in to another computer)
- Spark uses SSH to communicate between nodes (in an enterprise installation, these will run on many different computers)
- For the exercise we are going to set up SSH without a password

Check for Passphraseless SSH

- Type the command:
`$ ssh localhost`
- If you are prompted for a passphrase, you will need to set up passphraseless SSH

Set up Passphraseless SSH

- Type these commands:
`$ ssh-keygen -t dsa -P '' -f ~/.ssh/id_dsa`
`$ cat ~/.ssh/id_dsa.pub >> ~/.ssh/authorized_keys`
- If you don't do this, you will be prompted for your password whenever you start up / shut down Hadoop or Spark

Windows User Setup

Create a spa16 User in Windows

- You need to add a user with a name without spaces
- You can't do this from the Windows UI (since it demands a first and last name)
- Run a Windows Command Prompt as Administrator and enter the command:
`C:\Windows\System32>net user spa16 /add`
- Make this account an Administrator
 - Use Settings -> Accounts (Family & other users -> Change account type)
- Log out of Windows, and log back in again as the spa16 user

Software Prerequisites

Install Java (the full JDK, not the JRE) if not installed already

- <http://www.oracle.com/technetwork/java/javase/downloads/index.html>
- For Mac El Capitan, these commands will work if you have Homebrew:
\$ brew tap caskroom/cask
\$ brew unlink brew-cask
\$ brew cask install java
- For Debian and Ubuntu, these commands should work:
\$ sudo apt-get update
\$ sudo apt-get install default-jdk
- I'm not entirely sure about this, I already had java...

Install Python 2.7 if not installed already

- <https://www.python.org/downloads/>
- you don't need to know Python programming for this session, but need it to run some of the tools
- you may need to add C:\Python27 to %PATH% (see later slide for instructions)

Install git if not installed already

- <https://git-scm.com/downloads> (or apt, yum etc for Linux)
- Windows and Mac versions are on the USB stick in the downloads folder

Windows Software Prerequisites

Install the Visual C++ Runtime

- Run `vcredist_x64.exe` (it's on the on the USB stick in the downloads\windows directory)
- *I downloaded this from <https://www.microsoft.com/en-us/download/details.aspx?id=13523>*

Clone the Project Files

Start a Shell / Command Prompt

- OS X: Applications → Terminal; Linux: Start Terminal from the toolbar; Windows: run a Command Prompt

Create the Project Subdirectory

\$ `mkdir -p $HOME/SPA_2016/` ← for OS X and Linux

C:\> `mkdir C:\Users\spa16\SPA_2016` ← for Windows

Clone the Project files

\$ `cd $HOME/SPA_2016` ← for OS X and Linux

C:\> `cd \Users\spa16\SPA_2016` ← for Windows

\$ `git clone https://github.com/rozanski/bcs_spa16.git .`



Copy the Additional Directories from the USB Stick

- Copy the `datasets` and `download` directories from the USB stick into the `SPA_2016` directory

Install Hadoop

For OS X and Linux

- The Unix Hadoop software is in the `downloads/osx-linux` directory on the USB stick
- Extract `hadoop*.tar` into `~/SPA_2016/hadoop`

For Windows

- The standard Hadoop binaries do not work in Windows
- A custom Hadoop build for Windows is in the `downloads/windows` directory on the USB stick
- It includes Windows DLLs `hdfs.dll` and `hadoop.dll`, and batch scripts for starting the various programs (which generally don't work...)
- Extract `hadoop*.tar` into `C:\Users\spa16\SPA_2016\hadoop`

Check Extract

- Check you have directories `hadoop/bin`, `hadoop/sbin`, `hadoop/logs`, `hadoop/sbin...`

Install Spark and Hive

Install Spark

- The Spark software (all platforms) is in the `downloads/` directory on the USB stick
- Extract `spark*.tar` into `~/SPA_2016/spark` (OS X, Linux) or `C:\Users\spa16\SPA_2016\spark` (Windows)
- You should end up with directories `spark/bin`, `spark/conf`, `spark/logs`, `spark/sbin`...

Install Spark CSV Support

- This is also in the `downloads/` directory on the USB stick
- Copy `spark-csv*.jar` into `~/SPA_2016/spark/lib` (OSX and Linux) or `C:\Users\spa16\SPA_2016\spark\lib` (Windows)

Install Hive

- The Hive software (all platforms) is in the `downloads/` directory on the USB stick
- Extract `apache-hive*.tar` into `~/SPA_2016/hive` (OS X, Linux) or `C:\Users\spa16\SPA_2016\hive` (Windows)
- You should end up with directories `hive/bin`, `conf`...

Set Environment Variables (OS X, Linux)

Check the script `$HOME/SPA_2016/env.src`

- This attempts to derive `$JAVA_HOME` for your environment
 - It is configured for the latest Java version (1.8.0_91)
- It sets `$SPA_HOME` to the root directory for your project files
- It sets various environment variables for Hadoop, Spark and Hive

Run the Script

```
$ source $HOME/SPA_2016/env.src
```

- If there are no errors, and `$JAVA_HOME` and `$SPA_HOME` have been set correctly, you (probably) don't need to change it
- You will run the script at the start of each exercise

Set Environment Variables (Windows)

Set System Environment Variables

- You need Administrator access to do this
- Use Settings -> System Properties -> Advanced -> Environment Variables...
- You may need to restart

%JAVA_HOME%

- Set it to something like `C:\PROGRA~1\Java\JDK18~1.0_9`
- to get the 8.3 path, type `for %I in (.) do echo %~sI`

%PATH%

- Add the following paths to %PATH%:
`C:\WINDOWS\SYSTEM32`
`C:\Users\spa16\SPA_2016\hadoop\bin`
`C:\Users\spa16\SPA_2016\hadoop\sbin`

Run the Environment script

`C:\> env.cmd`

- You should run the script at the start of each exercise

Configure the Big Data Software

Sample Configuration Files

- Sample configuration files for the session are in directories in the `config` directory
- You will be copying the files into the software directories, and then confirming they are right for your setup

Copy Hadoop Config Files

- from: `config/hadoop_etc_hadoop`
- to: `hadoop/etc/hadoop`

Copy Spark Config Files

- from: `config/spark_conf`
- to: `spark/conf`

Copy Hive Config Files

- from: `config/hive_conf`
- to: `hive/conf`

On Ubuntu, you will also have to run the command:

```
sudo ln -s /home /Users
```

Review Hadoop Configuration

Hadoop Configuration Files

- Hadoop configuration files are XML files
- Check the paths and usernames shown below

`hadoop/etc/hadoop/core-site.xml`

<code>fs.defaultFS</code>	<code>hdfs://localhost:9000</code>	NameNode URI
<code>hadoop.tmp.dir</code>	<code>/Users/spa16/SPA_2016/data/hadoop/tmp</code>	Hadoop temporary directory
<code>hadoop.proxyuser.spa16.hosts</code>	<code>* (asterisk)</code>	used to configure connect connections; you adopt your operating system username and group when logged into Hadoop or Hive
<code>hadoop.proxyuser.spa16.groups</code>	<code>* (asterisk)</code>	

`hadoop/etc/hadoop/hdfs-site.xml`

<code>dfs.datanode.data.dir</code>	<code>file:/Users/spa16/SPA_2016/data/hadoop/hdfs/datanode</code>	Path on the local filesystem where the DataNode stores its blocks
<code>dfs.namenode.name.dir</code>	<code>file:/Users/spa16/SPA_2016/data/hadoop/hdfs/namenode</code>	Path on the local filesystem where the NameNode stores the namespace and transaction logs

Review Spark Configuration

Spark Configuration Files

- Edit Spark configuration script and files as shown below

`spark/conf/spark-env.sh`

<code>export HADOOP_CONF_DIR=/Users/spa16/SPA_2016/hadoop/etc/hadoop</code>	Tells Spark where the Hadoop configuration files can be found
<code>export SPARK_LOCAL_DIRS=/Users/spa16/SPA_2016/data/spark</code>	Tells Spark where to put its local storage
<code>export SPARK_WORKER_MEMORY=2000m</code>	how much total memory workers have to give executors

`spark/conf/slaves`

<code>localhost (or IP address)</code>	Tells Spark that a slave (Worker) will be running on the local computer
--	---

`spark/conf/spark-defaults.conf`

<code>spark.jars.packages com.databricks:spark-csv_2.11:1.4.0</code>	All on one line. Tells Spark that a slave (Worker) will be running on the local computer. Make sure the version (11.1.4.0) matches the version of the JAR you downloaded previously
--	---

Review Hive Configuration

Hive Configuration

- Edit Hive configuration files as shown below

`$SPA_2016/hive/conf/hive-site.xml`

<code>fs.defjvax.jdo.option. ConnectionURLaultFS</code>	<code>jdbc:derby;;databaseName=/Users/spa16/SPA_2016/ data/hive/metastore_db; create=true</code>	JDBC connect string for the metastore
<code>hive.execution.engine</code>	<code>spark</code>	Hive execution engine. Options are: mr (map reduce, default, deprecated), tez, spark.
<code>datanucleus.autoCreateTables</code>	<code>TRUE</code>	automatically create metadata tables on first invocation of hive server

`$SPA_2016/hive/conf/spark-defaults.conf`

<code>spark.master</code> spark://master:7077	Tells Hive the URL of the Spark Master server
<code>spark.jars.packages</code> <code>com.databricks:spark-csv_2.11:1.4.0</code>	Tells Hive to load Spark CSV support

Create Hive Metastore Database

- The Hive megastore database uses Derby to store Hive metadata
- In production, you would use a robust multi-threaded database like MySQL or SQL Server

OS X and Linux

- Run these commands

```
$ source $HOME/SPA_2016/env.src
```

```
$ spa_2016.bash init_metastore
```

- You should have a directory `$SPA_2016/data/hive/metastore_db/`

ALREADY CREATED!

Windows

- The above commands don't seem to work...

Real-World Big Data in Action

Hadoop Exercises

A Big Data Virtual
Filesystem



Initialise Hadoop Filesystem (OS X, Linux)

Format the HDFS Filesystem

- Equivalent to formatting an operating system filesystem partition
- Warning: this destroys all HDFS data!

```
$ source $HOME/SPA_2016/env.src
```

```
$ $HADOOP_PREFIX/bin/hdfs namenode -format
```

- Confirm there are no WARN or ERROR messages
- Do this before starting Hadoop
- If it goes wrong, delete the contents of the data directory before retrying

Check it has worked

```
$ ls $SPA_2016/data/hadoop/hdfs/  
namenode
```

- The directory `$SPA_2016/data/hadoop` stores the Hadoop physical operating system files for the namenode and datanode

Initialise Hadoop Filesystem (Windows)

Format the HDFS Filesystem

- Equivalent to formatting an operating system filesystem partition
- Warning: this destroys all HDFS data!

```
C:\> env.cmd
```

```
C:\> hadoop.cmd namenode -format
```

- Confirm there are no WARN or ERROR messages (ignore the DEPRECATED message)
- Do this before starting Hadoop

Check it has worked

```
C:\> dir C:\Users\spa16\SPA_2016\data\hadoop\hdfs  
namenode
```

- The directory `hadoop_data\hadoop` stores the Hadoop physical operating system files for the namenode and datanode



Start Hadoop Server (OS X, Linux)

Start Hadoop

```
$ source $HOME/SPA_2016/env.src
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start namenode
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start secondarynamenode
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start datanode
$ $HADOOP_PREFIX/sbin/yarn-daemon.sh start resourcemanager
$ $HADOOP_PREFIX/sbin/yarn-daemon.sh start nodemanager
```

- can also just run

```
$ $HADOOP_PREFIX/sbin/start-dfs.sh
$ $HADOOP_PREFIX/sbin/start-yarn.sh
```

To Stop Hadoop at any time

```
$ source $HOME/SPA_2016/env.src
$ $HADOOP_PREFIX/sbin/stop-dfs.sh
```

- Ignore messages like “Unable to load native-hadoop library for your platform”

Start Hadoop Server (Windows)

Start Hadoop

- Start a Command Prompt, then type
C:\> `start-dfs.cmd`
C:\> `start-yarn.cmd`
- Four command windows should open:
 - Hadoop namenode and datanode
 - YARN nodemanager and resource manager
- Check that the servers are all running and none terminated with errors
- If nothing starts, check %PATH% (should include `hadoop\bin` and `hadoop\sbin`)
- A nearly full disk may cause startup to fail



To Stop Hadoop at any time

- ```
C:\> stop-dfs.cmd
```
- ```
C:\> stop-yarn.cmd
```
- or just click on each window and type Ctrl-C

Check Hadoop is Running

Check Running Processes

`$ jps | sort -k 2` ← for OS X and Linux
`C:\> %JAVA_HOME%\bin\jps` ← for Windows

<code>nnnnn</code>	DataNode	← Hadoop datanode
<code>nnnnn</code>	NameNode	← Hadoop namenode
<code>nnnnn</code>	NodeManager	← Hadoop YARN node manager
<code>nnnnn</code>	ResourceManager	← Hadoop YARN resource manager
<code>nnnnn</code>	SecondaryNameNode	← Hadoop secondary namenode (may not have this)

Check Log Files

- Review log files in `logs` directory (OS X and Linux) or in the consoles (Windows)
- check there are no ERROR messages (a few WARN messages is usually ok)

Check Web Interfaces

- Hadoop Web UI <http://localhost:50070> (or <http://IPADDRESS:50070> for a VPS)
- try Utilities → Browse the Filesystem (it's empty at the moment)

Hadoop Command Line (OS X and Linux)

Hadoop Command Line

- many Unix shell file manipulation commands (`ls`, `mkdir`, `rm` etc) have Hadoop equivalents using `hadoop fs -<command>`
- for example: `$HADOOP_PREFIX/bin/hadoop fs -ls /user`
- see <https://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>

Create Your User Directories on the Hadoop Filesystem

```
$ source $SPA_2016/env.src
$ $HADOOP_PREFIX/bin/hadoop fs -mkdir -p /user/spa16/load/lfb
$ $HADOOP_PREFIX/bin/hadoop fs -mkdir -p /user/spa16/load/lhp
```

Check it's Worked

```
$ $HADOOP_PREFIX/bin/hadoop fs -ls /user/spa16/load
```

Hadoop Command Line (Windows)

Hadoop Command Line

- many Unix shell file manipulation commands (`ls`, `mkdir`, `rm` etc) have Hadoop equivalents using `%SPA_2016%\hadoop\bin\hdfs.cmd dfs -<command>`
- for example: `%SPA_2016%\hadoop\bin\hdfs.cmd -ls /user`
- see <https://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>
- Note this script is in `hadoop/bin`, not `hadoop/sbin`

Create Your User Directories on the Hadoop Filesystem

```
C:\> hdfs.cmd dfs -mkdir -p /user/spa16/load/lfb
```

```
C:\> hdfs.cmd dfs -mkdir -p /user/spa16/load/lhp
```

Check it's Worked

```
C:\> hdfs.cmd dfs -ls /user/spa16/load
```

Found 2 items

drwxr-xr-x	-	spa16	supergroup	0	2016-06-19	15:16	/user/spa16/load/lfb
drwxr-xr-x	-	spa16	supergroup	0	2016-06-19	15:16	/user/spa16/load/lhp

Load Some Data into Hadoop

London Fire Brigade Reported Incidents

- Original from <http://data.london.gov.uk/dataset/london-fire-brigade-incident-records>
- Covers the period 2013 - 2016
- I loaded it into Excel and converted into a 'Windows Comma-Separated' file
- You will find it in:
 - `$SPA_2016/datasets/lfb/load/lfb.csv` (*OS X, Linux*)
 - `C:\Users\spa16\SPA_2016\datasets\lfb\load\lfb.csv` (*Windows*)
- There is a larger file, `lfb-large.csv`, if you want to explore with more data

Load the data into Hadoop

- Run the following command (split over multiple lines here for readability, enter all one one line):
- for OS X and Linux:

```
$ $HADOOP_HOME/bin/hadoop fs -put  
  $SPA_2016/datasets/lfb/load/lfb.csv  
  hdfs://localhost:9000/user/spa16/load/lfb
```

- for Windows:

```
C:\> hdfs.cmd dfs -put  
  C:\Users\spa16\SPA_2016\datasets\lfb\load\lfb.csv  
  hdfs://localhost:9000/user/spa16/load/lfb
```

Check It Has Loaded Into Hadoop

Browse Hadoop from the Command Line

```
$ $HADOOP_HOME/bin/hadoop fs -ls /user/spa16/load/lfb
```

```
C:\> hdfs.cmd dfs -ls /user/spa16/load/lfb
```

```
Found 1 items
```

```
-rw-r--r--    3 spa16 supergroup    79888721 2016-05-29 10:53 /user/spa16/load/lfb.csv
```

- count the number of lines in the file

```
$ $HADOOP_HOME/bin/hadoop fs -cat /user/spa16/load/lfb/lfb.csv | wc -l
```

```
C:\> hadoop.cmd fs -cat /user/spa16/load/lfb/lfb.csv | find /c /v ""
```

```
322217
```

```
322218 (for Windows)
```

Browse Hadoop from your Web Browser

- <http://localhost:50070/explorer.html>
- look in /user/spa16/load/lfb

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	spa16	supergroup	76.19 MB	6/5/2016, 5:46:50 PM	3	128 MB	lfb.csv

Real-World Big Data in Action

Spark Exercises

A Big Data Processing
Engine



Start Spark Server (OS X and Linux)

Start Hadoop if not already running

- see previous slides (for Windows, don't forget to use a Command Prompt)

Start Spark Server

- run the following commands

```
$ source $HOME/SPA_2016/env.src
```

```
$ $SPARK_HOME/sbin/start-master.sh
```

```
$ $SPARK_HOME/sbin/start-slaves.sh spark://ipaddress:7077
```

- `ipaddress` is the IP address of your computer (hostname may work)
- enter your password if prompted (Spark uses `ssh`)

Stop Spark Server

```
$ source $HOME/SPA_2016/env.src
```

```
$ $SPARK_HOME/sbin/stop-all.sh
```

Start Spark Server (Windows)

Start Hadoop if not already running

- see previous slides (for Windows, don't forget to use a Command Prompt)

Start Spark Server

- use a Command Prompt

```
C:\> cd \Users\spa16\SPA_2016
```

```
C:\> start spark\bin\spark-class.cmd \
    org.apache.spark.deploy.master.Master
```

```
C:\> start spark\bin\spark-class.cmd \
    org.apache.spark.deploy.worker.Worker spark://IPADDRESS:7077
```

- **IPADDRESS** is the IP address of your computer (it's logged by the Master)
- the Worker may need several attempts to connect to the master

Stop Spark Server

- press Control-C in each Spark window or click the red X in the corner



Check Hadoop and Spark Are Running

Check Running Processes

<code>\$ jps sort -k 2</code>	← for OS X and Linux
<code>C:\> %JAVA_HOME%\bin\jps</code>	← for Windows
<code>nnnnn DataNode</code>	← Hadoop datanode
<code>nnnnn Master</code>	← Spark master
<code>nnnnn NameNode</code>	← Hadoop namenode
<code>nnnnn NodeManager</code>	← Hadoop YARN node manager
<code>nnnnn ResourceManager</code>	← Hadoop YARN resource manager
<code>nnnnn SecondaryNameNode</code>	← Hadoop secondary namenode
<code>nnnnn Worker</code>	← Spark slave

Check Log Files

- Review log files in `logs` directory (OS X and Linux) or in the consoles (Windows)
- check there are no ERROR messages (a few WARN messages is usually ok)

Check Web Interfaces

- Hadoop Web UI <http://localhost:50070>
- Browse Hadoop Filesystem <http://localhost:50070/explorer.html#>
- Spark Web UI <http://localhost:8080>

Pyspark

Pyspark

- Pyspark allows you to submit Spark commands from a Python shell, in the same way you would invoke Spark programatically
- Pyspark is a wrapper script for `spark-submit`, which is a script you use to launch Spark applications (jar files) on a Spark cluster

Launching Pyspark

- Start Hadoop and Spark
- Start Pyspark:

```
$ $SPARK_HOME/bin/pyspark --master spark://IPADDRESS:7077
```

```
C:\> \SPA_2016\env.cmd
```

```
C:\> %SPARK_HOME%\bin\pyspark.cmd --master spark://IPADDRESS:7077
```

- If you get connection refused, try `localhost` or your computer's hostname (`spa16` for the Digital Ocean VPS)

- You should get the message:

```
SparkContext available as sc, HiveContext available as sqlContext.
```

- You can run any Python command at this point
 - You can also call functions in the `pyspark.sql` library
- ```
>>> help(sqlContext)
```



---

# Data Science Using Spark (1 of 2)

---

## Load the LFB Data from the Hadoop Filesystem into Spark

- Enter the following command at the Pyspark prompt (*on one line, split here for readability*)

```
>>> lfb = sqlContext.read.format("com.databricks.spark.csv").
 option("header", "true").option("inferSchema", "true").
 option("mode", "DROPMALFORMED").
 load("hdfs://localhost:9000/user/spa16/load/lfb/lfb.csv")
```

## Check it's Loaded

```
>>> print lfb.count()
...
322217
```

## Display the data column names

```
>>> lfb.printSchema()
```

## Look at Some Data

```
>>> lfb.filter(lfb.IncidentGroup == "Special Service").limit(5).show()
```



---

# Data Science Using Spark (2 of 2)

---

## Incident Counts by Type

```
>>> lfb.groupBy("IncidentGroup").count().show()
```

## Incident Counts by Stop Code

```
>>> lfb.groupBy("StopCodeDescription").count().show(truncate=False)
```

## Most Dangerous Areas

```
>>> lfb.groupBy("Postcode_district").count(). \
 sort("count", ascending=False).show()
```

## And What Happens There

```
>>> lfb.rollup("IncidentGroup", "Postcode_district"). \
 count().sort("count", ascending=False).show()
```

## “Frequent” Problem Areas

```
>>> for borough in sorted(lfb.freqItems(["IncGeo_BoroughName"]).first()[0]):
 print borough ← this line starts with a tab or some spaces (this is Python!)
```

*Real-World Big Data in Action*

---

# Hive Exercises

A Big Data data warehousing  
infrastructure





---

# Start Hive Server (OS X and Linux)

---

## Start Hadoop and Spark if not already running

- see earlier slides

## Start Hive Server

```
$ source $HOME/SPA_2016/env.src
$ nohup $HIVE_HOME/bin/hive --service hiveserver2 2>&1 > /dev/null &
```

## Stop Hive Server

```
$ killall HiveServer2
```

## You can use my Start / Stop Script to start all servers

```
$ source $HOME/SPA_2016/env.src
$ spa_2016.bash start
...
$ spa_2016.bash stop
```

---

# Start Hive Server (Windows)

---

## Start Hadoop and Spark if not already running

- see earlier slides

## Start Hive Server

```
C:\> cd \SPA_2016
C:\> env.cmd
C:\> START %HIVE_HOME%\bin\hive.cmd --service hiveserver2
```

## Stop Hive Server

- press Control-C in the Hive window or click the red X in the corner

## You can use my Start / Stop Script to start all servers

```
C:\> cd \SPA_2016
C:\> spa_2016.cmd start
...
C:\> spa_2016.cmd stop
```





---

# Beeline

---

## Beeline

- Beeline allows you to run Hive SQL queries from a command shell
- Beeline commands can span multiple lines and are terminated by a semicolon ;
- Exit Beeline by typing `!quit` at the prompt

## Run Beeline

```
$ $SPARK_HOME/bin/beeline -u jdbc:hive2:// --color
```

- *Do not run the version of Beeline in \$HIVE\_HOME/bin!*

```
C:\> cd %SPA_2016%
```

```
C:\> spa_2016.cmd beeline
```

## Check your Hive databases

```
0: jdbc:hive2://> SHOW DATABASES;
```

```
+-----+
| database_name |
+-----+
| default |
+-----+
```

- You have an empty Hive installation

---

# Data Science Using Hive (1 of 2)

---

## Create your database

- Start beeline and enter the command:

```
0: jdbc:hive2://> create database spa_2016;
0: jdbc:hive2://> show databases;
```

## Load the LFB Data into Hive

- Quit beeline (!quit) and run this script, which creates a Hive external table called `lfb_data`

```
$ $SPARK_HOME/bin/beeline -u jdbc:hive2:// --color < \
 $SPA_2016/datasets/lfb/load_external.hive
C:\> spa_2016.cmd beeline < C:\Users\spa16\SPA_2016\datasets\lfb\load_external.hive
```

## Check It's Loaded

- Run these commands in beeline

```
0: jdbc:hive2://> use spa_2016;
0: jdbc:hive2://> describe lfb_data;
0: jdbc:hive2://> select count(*) from lfb_data;
```

- The table should contain 322,217 rows



---

# Data Science Using Hive (2 of 2)

---

## Incident Counts by Type

```
//> select incidentgroup, count(*) from lfb_data group by incidentgroup;
```

## Incident Counts by Stop Code

```
//> select stopcodedescription, count(*) from lfb_data
 group by stopcodedescription;
```

## Most Dangerous Areas

```
//> select postcode_district, incgeo_boroughname, count(*) as c
 from lfb_data group by postcode_district, incgeo_boroughname
 having c > 1000 order by c desc limit 10;
```

## And What Happens There

```
//> select postcode_district, incidentgroup, count(*) as c
 from lfb_data group by postcode_district, incidentgroup
 having c > 1000 order by c desc;
```

## “Frequent” Problem Areas

- There is no Hive equivalent to Spark `freqItems`

*Real-World Big Data in Action*

---

# Additional Exercises

**Even Bigger Big Data**

---



---

# Load Even Bigger Data into Hadoop

---

## London House Prices

- Original from <http://data.london.gov.uk/dataset/average-house-prices-borough>
- Four CSV files, covers the period 1995 - 2014
- You can find them in `$SPA_2016/datasets/lfb/load/*.csv`

## Load the data into Hadoop using `hadoop fs`

```
$ $HADOOP_HOME/bin/hadoop fs -put \
 $SPA_2016/datasets/lfb/load/*.csv \
 hdfs://localhost:9000/user/spa16/load/lhp
```

## Load the data from Hadoop into Spark using Pyspark

```
>>> lfb = sqlContext.read.format('com.databricks.spark.csv').
 option('header', 'true').option('inferSchema', 'true').
 option('mode', 'DROPMALFORMED').load('/user/spa16/load/lhp/*.csv')
```

## Load the data from Hadoop into Hive using Spark Beeline

- Run this script, which creates a Hive external table called `lfb_data`
- ```
$ $SPARK_HOME/bin/beeline -u jdbc:hive2:// --color < \  
  $SPA_2016/datasets/lfb/load_external.hive
```

Data Science With Bigger Data

Average London House Prices by Year

```
//> select year, avg(price) avg_price from lhp_data  
      group by year order by year;
```

Average Borough House Prices by Year

```
//> select local_authority, year, avg(price) avg_price from lhp_data  
      group by local_authority, year order by local_authority, year;
```

Maximum and Minimum House Prices by Year

```
//> select year, max(price) max_price, min(price) min_price from lhp_data  
      group by year order by year;
```

Prices by Property Type and Year

```
//> select property_type, year, avg(price) avg_price from lhp_data  
      group by property_type, year order by avg_price desc;
```

Prices by Tenure and Year

```
//> select tenure, year, count(*) count, avg(price) avg_price from lhp_data  
      group by tenure, year order by tenure, year;
```

- I assume tenure means freehold or leasehold etc.

Real-World Big Data in Action

Appendix

**Further Information and
Troubleshooting**

Web UIs

Hadoop

- Namenode: <http://localhost:50070/>
- Datanodes: <http://localhost:50075/>
- Secondary Namenode: <http://localhost:50090/>
- YARN Resource Manager: <http://localhost:8088/cluster>
- see: <http://blog.cloudera.com/blog/2009/08/hadoop-default-ports-quick-reference/>

Spark

- Spark management interface: <http://localhost:8080>
- PySpark UI: <http://localhost:4040/>

Hive

- as defined in hive-site.xml

Troubleshooting

Hints and Tips

- logging out or rebooting in Windows fixes many problems
- avoid pathnames with spaces - this will break many commands
- don't forget to `source env.src` (OS X / Linux) or run `env.bat` (Windows)
- on Windows, make sure that `hadoop\bin` and `%JAVA_HOME%\bin` are in your `%PATH%`
- on Windows you may need to clear the contents of the temporary directory
`C:\Users\spa16\AppData\Local\Temp`

Install Cygwin (Windows 64-bit only)

- Provides a BASH shell to run scripts and commands (you can't use it to run the Big Dat tools though)
- Download from <https://cygwin.com/install.html>
- A bit easier to use for debugging than the Windows Command Prompt

Troubleshooting Commands

- what is listening on a port?
\$ `sudo lsof -i -n -P | grep TCP | grep $PORT` # OS X
\$ `sudo netstat -tulpn | grep :$PORT` # Linux
- set Hadoop debug level
\$ `$HADOOP_HOME/bin/hadoop daemonlog -setlevel 127.0.0.1:50070 \`
`org.apache.hadoop.hdfs.server.namenode.NameNode DEBUG`

LFB Hive Schema

root

```
|-- IncidentNumber: string (nullable = true)
|-- DateOfCall: string (nullable = true)
|-- TimeOfCall: string (nullable = true)
|-- IncidentGroup: string (nullable = true)
|-- StopCodeDescription: string (nullable = true)
|-- SpecialServiceType: string (nullable = true)
|-- PropertyCategory: string (nullable = true)
|-- PropertyType: string (nullable = true)
|-- AddressQualifier: string (nullable = true)
|-- Postcode_full: string (nullable = true)
|-- Postcode_district: string (nullable = true)
|-- IncGeo_BoroughCode: string (nullable = true)
|-- IncGeo_BoroughName: string (nullable = true)
|-- IncGeo_WardCode: string (nullable = true)
|-- IncGeo_WardName: string (nullable = true)
|-- Easting_m: string (nullable = true)
|-- Northing_m: string (nullable = true)
|-- Easting_rounded: integer (nullable = true)
|-- Northing_rounded: integer (nullable = true)
|-- FRS: string (nullable = true)
|-- IncidentStationGround: string (nullable = true)
|-- FirstPumpArriving_AttendanceTime: string (nullable = true)
|-- FirstPumpArriving_DeployedFromStation: string (nullable = true)
|-- SecondPumpArriving_AttendanceTime: string (nullable = true)
|-- SecondPumpArriving_DeployedFromStation: string (nullable = true)
|-- NumStationsWithPumpsAttending: string (nullable = true)
|-- NumPumpsAttending: string (nullable = true)
```

LHP Hive Schema

```
root
|-- id: string (nullable = true)
|-- transaction_id: string (nullable = true)
|-- price: string (nullable = true)
|-- date_processed: string (nullable = true)
|-- quarter: string (nullable = true)
|-- month: string (nullable = true)
|-- year: string (nullable = true)
|-- year_month: string (nullable = true)
|-- post_code: string (nullable = true)
|-- property_type: string (nullable = true)
|-- whether_new: string (nullable = true)
|-- tenure: string (nullable = true)
|-- address1: string (nullable = true)
|-- address2: string (nullable = true)
|-- address3: string (nullable = true)
|-- address4: string (nullable = true)
|-- town: string (nullable = true)
|-- local_authority: string (nullable = true)
|-- county: string (nullable = true)
|-- record_status: string (nullable = true)
|-- post_code_clean: string (nullable = true)
|-- inner_outer: string (nullable = true)
|-- borough_code: string (nullable = true)
|-- borough_name: string (nullable = true)
|-- ward_code: string (nullable = true)
|-- ward_name: string (nullable = true)
|-- msoa11: string (nullable = true)
|-- lsoa11: string (nullable = true)
|-- oall: string (nullable = true)
```