

SPA Conference 2016

Real-World Big Data in Action

Nick Rozanski
Eoin Woods
Chris Cooper-Bland



Number 2 in an occasional series

Prerequisites

Install Java Runtime (JDK)

- <http://www.oracle.com/technetwork/java/javase/downloads/index.html>
- For Mac El Capitan, these commands might also work:
`$ brew tap caskroom/cask`
`$ brew unlink brew-cask`
`$ brew cask install java`
- For Ubuntu 16, these commands might also work:
`$ sudo apt-get update`
`$ sudo apt-get install default-jdk`
- I'm not entirely sure about this, I already had the JRE...

Install Python 2.7

- <https://www.python.org/downloads/>
- you don't need to know Python programming for this session, but need it to run some of the tools
- Python is probably installed already

Create Your Project Directory

Start a BASH Shell

- OS X: Applications → Terminal
- Linux: Start Terminal from the toolbar

Create the Project Subdirectory

```
$ mkdir -p $HOME/SPA_2016/
```

Clone the Project files

```
$ cd $HOME/SPA_2016
```

```
$ git clone https://github.com/rozanski/bcs\_spa16.git .
```

- don't miss out the dot at the end of the command!
- otherwise you will have to:

```
mv $HOME/SPA_2016/bcs_spa16/* $HOME/SPA_2016/
```



Set Environment Variables

Check the script `$HOME/SPA_2016/env.src`

- This attempts to derive `$JAVA_HOME` for your environment
 - It is configured for the latest Java version (1.8.0_91)
- It sets `$SPA_HOME` to the root directory for your project files
- It sets various environment variables for Hadoop, Spark and Hive

Run the script

```
$ source $HOME/SPA_2016/env.src
```

- If there are no errors, and `$JAVA_HOME` and `$SPA_HOME` have been set correctly, you (probably) don't need to change it...

Install the Big Data Software

Download Hadoop into \$SPA_2016/hadoop

- <https://www.apache.org/dyn/closer.cgi/hadoop/common/>
- Select stable source, download and extract the binary tarball
- You should end up with directories `spark/bin`, etc, `logs`, `sbin`...
- (you could also use `brew` for Mac, `apt` for Linux, but we won't for this demo)

Download Spark into \$SPA_2016/spark

- <https://spark.apache.org/downloads.html>
- Choose the package type “pre-built for Hadoop 2.6 and later”
- Download and extract the binary tarball
- You should end up with directories `hadoop/bin`, `conf`, `logs`, `sbin`...

Download Hive into \$SPA_2016/hive

- <https://www.apache.org/dyn/closer.cgi/hive/>
- Download and extract the latest binary tarball
- You should end up with directories `bin`, `conf`...

*I will also provide
the software on a
USB stick*

Real-World Big Data in Action

Hadoop

A Big Data Virtual
Filesystem



Configure Hadoop Core Settings

Edit `$HADOOP_PREFIX/etc/hadoop/core-site.xml`

- used by the Hadoop client to access the Hadoop filesystem
- a template can be found in directory `$SPA_2016/config`
- add the following lines:

```
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
  <description>NameNode URI</description>
</property>
<property>
  <name>hadoop.proxyuser.YOURNAME.hosts</name>
  <value>*</value>
</property>
<property>
  <name>hadoop.proxyuser.YOURNAME.groups</name>
  <value>*</value>
</property>
```

- replace `YOURNAME` with your operating system user name (no spaces!)

Configure Hadoop Site Settings

Edit \$HADOOP_PREFIX/etc/hadoop/hdfs-site.xml

- tells Hadoop where to put the physical operating system files for the datanode and namenode
- a template can be found in directory \$SPA_2016/config
- change /YOURHOME to your home directory (eg /home/nick or /Users/nick)

```
<property>
```

```
  <name>dfs.datanode.data.dir</name>
```

```
  <value>file:///YOURHOME/SPA_2016/data/hadoop/hdfs/datanode</value>
```

```
  <description>
```

```
    Paths on the local filesystem where the DataNode stores its  
blocks
```

```
  </description>
```

```
</property>
```

```
<property>
```

```
  <name>dfs.namenode.name.dir</name>
```

```
  <value>file:///YOURHOME/SPA_2016/data/hadoop/hdfs/namenode</value>
```

```
  <description>
```

```
    Path on the local filesystem where the NameNode stores the  
namespace and transaction logs
```

```
  </description>
```

```
</property>
```

Initialise Hadoop Filesystem

Format the HDFS Filesystem

- Equivalent to formatting an operating system filesystem partition
- Warning: this destroys all HDFS data!

```
$ source $HOME/SPA_2016/env.src
```

```
$ $HADOOP_PREFIX/bin/hdfs namenode -format
```

- Confirm there are no **WARN** or **ERROR** messages
- Do this before starting Hadoop

Check it has worked

```
$ ls $SPA_2016/data/hadoop/hdfs/  
namenode
```

- The directory `$SPA_2016/data/hadoop` directory will store the Hadoop physical operating system files

Start (and Stop) Hadoop Server

Start Hadoop

```
$ source $HOME/SPA_2016/env.src
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start namenode
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start secondarynamenode
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start datanode
# can also do $HADOOP_PREFIX/sbin/start-dfs.sh
```

To Stop Hadoop at any time

```
$ source $HOME/SPA_2016/env.src
$ $HADOOP_PREFIX/sbin/stop-dfs.sh
```

- You may be asked for your password (the scripts use SSH)
- Ignore messages like “Unable to load native-hadoop library for your platform”

Check Hadoop is Running

Check Running Processes

```
$ jps | sort -k 2
```

```
nnnnn DataNode
```

← Hadoop datanode

```
nnnnn NameNode
```

← Hadoop namenode

```
nnnnn SecondaryNameNode
```

← Hadoop secondary namenode

Check Log Files

```
egrep 'WARN|ERROR' $SPA_2016/hadoop/logs/hadoop-<user>-namenode-<hostname>.log
```

```
egrep 'WARN|ERROR' $SPA_2016/hadoop/logs/hadoop-<user>-datanode-<hostname>.log
```

```
egrep 'WARN|ERROR' $SPA_2016/hadoop/logs/hadoop-<user>-secondarynamenode-  
    <hostname>.log
```

- check there are no ERROR messages (a few WARN messages is usually ok)

Check Web Interfaces

- Hadoop Web UI <http://localhost:50070>
- try Utilities → Browse the Filesystem (it's empty at the moment)

Hadoop Command Line

Hadoop Command Line

- many Unix shell file manipulation commands (`ls`, `mkdir`, `rm` etc) have Hadoop equivalents using `hadoop fs -<command>`
- for example: `$HADOOP_PREFIX/bin/hadoop fs -ls /user`
- see <https://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>

Create Your User Directories on the Hadoop Filesystem

```
$ source $SPA_2016/env.src
$ $HADOOP_PREFIX/bin/hadoop fs -mkdir -p /user/YOURNAME/load/lfb
$ $HADOOP_PREFIX/bin/hadoop fs -mkdir -p /user/YOURNAME/load/lhp
# user identities map 1-1 from the O/S
```

- replace `YOURNAME` with your operating system user name

Check it's Worked

```
$ $HADOOP_PREFIX/bin/hadoop fs -ls /user/YOURNAME/load
```

Load Some Data into Hadoop

London Fire Brigade Reported Incidents

- Original from <http://data.london.gov.uk/dataset/london-fire-brigade-incident-records>
- Covers the period 2013 - 2016
- I loaded it into Excel and converted into a 'Windows Comma-Separated' file
- You can find it in `$SPA_2016/datasets/LFB/load/LFB.csv`
- There is a larger file, `LFB-large.csv`, if you want to play around with more data

Load the data into Hadoop

- Run the following command (split over three lines here for readability):

```
$ $HADOOP_HOME/bin/hadoop fs -put \  
    $SPA_2016/datasets/LFB/load/LFB.csv \  
    hdfs://localhost:9000/user/YOURNAME/load/lfb
```

Check It Has Loaded Into Hadoop

Browse Hadoop from the Command Line

```
$ $HADOOP_HOME/bin/hadoop fs -ls /user/YOURNAME/load/lfb
```

Found 1 items

```
-rw-r--r--  3 nick supergroup  79888721 2016-05-29 10:53 /user/nick/load/LFB.csv
```

- count the number of lines in the file

```
$ hadoop fs -cat /user/YOURNAME/load/lfb/LFB.csv | wc -l
```

Browse Hadoop from your Web Browser

- <http://localhost:50070/explorer.html>
- look in /user/YOURNAME/load/lfb

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	nick	supergroup	76.19 MB	6/5/2016, 5:46:50 PM	3	128 MB	LFB.csv

Real-World Big Data in Action

Spark

A Big Data Processing
Engine



Configure Spark

Edit `$SPA_2016/spark/conf/spark-env.sh`

- add the lines:
`export HADOOP_CONF_DIR=/YOURHOME/SPA_2016/hadoop/etc/hadoop`
`export SPARK_LOCAL_DIRS=/YOURHOME/SPA_2016/data/spark`
- where `YOURHOME` is your home directory
- these tell Spark where to find files on the local filesystem

Create `$SPA_2016/spark/conf/slaves`

- make sure the file includes:
`localhost`

Download Spark CSV Support

- Download `spark-csv` from <https://spark-packages.org/package/databricks/spark-csv>
- Save the latest JAR into `$SPA_2016/spark/lib`

Edit `$SPA_2016/spark/conf/spark-defaults.conf`

- add the line:
`spark.jars.packages com.databricks:spark-csv_2.11:1.4.0`
- make sure the version (11.1.4.0) matches the version of the JAR you downloaded!

Set up Passphraseless SSH

Spark and SSH

- SSH is Secure Shell, a cryptographically secure way of running services over an insecure network (for example, logging in to another computer)
- Spark uses SSH to communicate between nodes (in an enterprise installation, these will run on many different computers)
- For the exercise we are going to set up SSH without a password

Check for Passphraseless SSH

- Type the command:
`$ ssh localhost`
- If you are prompted for a passphrase, you will need to set up passphraseless SSH

Set up Passphraseless SSH

- Type these commands:
`$ ssh-keygen -t dsa -P '' -f ~/.ssh/id_dsa`
`$ cat ~/.ssh/id_dsa.pub >> ~/.ssh/authorized_keys`
- If you don't do this, you will be prompted for your password whenever you start up / shut down Hadoop or Spark

Start (and Stop) Spark Server

Start Hadoop if not already running

```
$ source $HOME/SPA_2016/env.src
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start namenode
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start secondarynamenode
$ $HADOOP_PREFIX/sbin/hadoop-daemon.sh start datanode
# can also do $HADOOP_PREFIX/sbin/start-dfs.sh
```

Start Spark Server

```
$ source $HOME/SPA_2016/env.src
$ $SPARK_HOME/sbin/start-master.sh
$ $SPARK_HOME/sbin/start-slaves.sh spark://hostname:7077
```

- where `hostname` is the host name (or IP address) of your computer
- enter your password if prompted (Spark uses `ssh`)

Stop Spark Server

```
$ source $HOME/SPA_2016/env.src
$ $SPARK_HOME/sbin/stop-all.sh
```

Check Hadoop and Spark Are Running

Check Running Processes

```
$ jps | sort -k 2
```

<i>nnnnn</i> DataNode	← Hadoop datanode
<i>nnnnn</i> Master	← Spark master
<i>nnnnn</i> NameNode	← Hadoop namenode
<i>nnnnn</i> SecondaryNameNode	← Hadoop secondary namenode
<i>nnnnn</i> Worker	← Spark slave

Check Log Files

- `egrep 'WARN|ERROR' $SPA_2016/spark/logs/*Master*.out`
- `egrep 'WARN|ERROR' $SPA_2016/spark/logs/*Worker*.out`

Check Web Interfaces

- Hadoop Web UI <http://localhost:50070>
- Browse Hadoop Filesystem <http://localhost:50070/explorer.html#>
- Spark Web UI <http://localhost:8080>

Pyspark

Pyspark

- Pyspark allows you to submit Spark commands from a Python shell, in the same way you would invoke Spark programatically
- Pyspark is a wrapper script for spark-submit, which is a script you use to launch Spark applications on a Spark cluster

Launching Pyspark

- Start Hadoop and Spark
- Start Pyspark:

```
$ $SPARK_HOME/bin/pyspark
```

- You should get the message:

```
SparkContext available as sc, HiveContext available as  
sqlContext.
```

- You can run any Python command at this point
- You can also call functions in the pyspark.sql library

```
$ help(sqlContext)
```

Let's Do Some Data Science!

Load the LFB Data into Spark

- Enter the following command at the Pyspark prompt (*on one line, split here for readability*)

```
>>> lfb = sqlContext.read.format('com.databricks.spark.csv').  
      option('header', 'true').option('inferSchema', 'true').  
      option('mode', 'DROPMALFORMED').load('/user/YOURNAME/load/lfb/LFB.csv')
```

Check it's Loaded

```
>>> print lfb.count()
```

```
...
```

```
322217
```

Display the data column names

```
>>> lfb.printSchema()
```

Look at Some Data

```
>>> lfb.filter(lfb.IncidentGroup == 'Special Service').limit(5).show()
```

Data Science continued

Incident Counts by Type

```
>>> lfb.groupBy('IncidentGroup').count().show()
```

Incident Counts by Stop Code

```
>>> lfb.groupBy('StopCodeDescription').count().show(truncate=False)
```

Most Dangerous Areas

```
>>> lfb.groupBy('Postcode_district').count(). \
    sort('count', ascending=False).show()
```

And What Happens There

```
>>> lfb.rollup('IncidentGroup', 'Postcode_district'). \
    count().sort('count', ascending=False).show()
```

“Frequent” Problem Areas

```
>>> for borough in sorted(lfb.freqItems(['IncGeo_BoroughName']).first()[0]):
    print borough ← this line starts with a tab or some spaces (this is Python!)
```

Real-World Big Data in Action

Hive

A Big Data data warehousing
infrastructure



Configure Hive (1 of 2)

Edit \$SPA_2016/hive/conf/hive-site.xml

```
<configuration>
  <property>
    <name>javax.jdo.option.ConnectionURL</name>
    <value>jdbc:derby::databaseName=/Users/YOURNAME/SPA_2016/data/hive/
metastore_db;create=true</value>
    <description>JDBC connect string for a JDBC metastore</description>
  </property>
  <property>
    <name>hive.execution.engine</name>
    <value>spark</value>
    <description>
      Expects one of [mr, tez, spark].
      Chooses execution engine. Options are: mr (Map reduce, default), tez, spark.
      While MR
      remains the default engine for historical reasons, it is itself a historical
      engine
      and is deprecated in Hive 2 line. It may be removed without further warning.
    </description>
  </property>
</configuration>
```

Configure Hive (2 of 2)

Edit `$SPA_2016/hive/conf/spark-defaults.conf`

- Ensure you have this line, which tells Hive where to find Spark
`spark.master spark://master:7077`
- and this line:
`spark.jars.packages com.databricks:spark-csv_2.11:1.4.0`

Create Metastore Database

- ```
$ source $HOME/SPA_2016/env.src
$ mkdir $SPA_2016/data/hive
$ cd $SPA_2016/data/hive
$ $HIVE_HOME/bin/schematool -initSchema -dbType derby
```
- You should have a directory `$SPA_2016/data/hive/metastore_db$`
  - (this may not be necessary!)

---

# Start (and Stop) Hive Server

---

## Start Hadoop if not already running

- see earlier slide

## Start Spark Server if not already running

- see earlier slide

## Start Hive Server

```
$ source $HOME/SPA_2016/env.src
```

```
$ nohup $HIVE_HOME/bin/hive --service hiveserver2 2>&1 > /dev/null &
```

## Stop Hive Server

```
$ killall HiveServer2
```



---

# Beeline

---

- Beeline allows you to run Hive SQL queries from a command shell
- Start Hadoop, Spark and Hive if not already running
- Start Beeline:

```
$SPARK_HOME/bin/beeline -u jdbc:hive2:// --color
```

- *Do not run the version of Beeline in \$HIVE\_HOME/bin!*
- Beeline commands can span multiple lines and are terminated by a semicolon ;

- Check your Hive databases

```
0: jdbc:hive2://> SHOW DATABASES;
```

```
+-----+---+
| database_name |
+-----+---+
| default |
+-----+---+
```

- You have an empty Hive installation
- Exit Beeline by typing `!quit` at the prompt

---

# Data Science Using Hive

---

## Create your database

- Start beeline and enter the command:

```
0: jdbc:hive2://> create database spa_2016;
```

## Load the LFB Data into Hive

- Run this script, which creates a Hive external table called `lfb_data`

```
$SPARK_HOME/bin/beeline -u jdbc:hive2:// --color < \
 datasets/LFB/load_external.hive
```

## Check It's Loaded

- Run these commands in beeline

```
$SPARK_HOME/bin/beeline -u jdbc:hive2://
0: jdbc:hive2://> use spa_2016;
0: jdbc:hive2://> select count(*) from lfb_data;
0: jdbc:hive2://> describe lfb_data;
```

- The table should contain 322,217 rows

---

# Data Science Continued

---

## Incident Counts by Type

```
//> select incidentgroup, count(*) from lfb_data group by incidentgroup;
```

## Incident Counts by Stop Code

```
//> select stopcodedescription, count(*) from lfb_data
 group by stopcodedescription;
```

## Most Dangerous Areas

```
//> select postcode_district, incgeo_boroughname, count(*) as c
 from lfb_data group by postcode_district, incgeo_boroughname
 having c > 1000 order by c desc limit 10;
```

## And What Happens There

```
//> select postcode_district, incidentgroup, count(*) as c
 from lfb_data group by postcode_district, incidentgroup
 having c > 1000 order by c desc;
```

## “Frequent” Problem Areas

- no Hive equivalent to Spark `freqItems`



*Real-World Big Data in Action*

---

Next Steps

**Some Other Big Data Tools**

---

---

# Next Steps

---

- different Hadoop configurations
- Cloudera VM
- Multiple slaves
- YARN



TO DO

*Real-World Big Data in Action*

---

# Appendix

**Further Information and  
Troubleshooting**

---



---

# Useful Links

---

## Hadoop FAQ

- <http://wiki.apache.org/hadoop/FAQ>

## Hadoop Filesystem commands Reference

- <http://hadoop.apache.org/docs/current/hadoop-project-dist/hadoop-common/FileSystemShell.html>

## Pyspark sqlContext Reference

- <https://spark.apache.org/docs/latest/api/python/pyspark.sql.html#pyspark.sql.DataFrame>

## Hive SQL Reference

- <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+DDL>
- <https://cwiki.apache.org/confluence/display/Hive/LanguageManual+DML>

---

# Troubleshooting

---

## Troubleshooting Commands

- what is listening on a port?

```
$ sudo lsof -i -n -P | grep TCP | grep $PORT # OS X
```

```
$ sudo netstat -tulpn | grep :$PORT # Linux
```

- set debug level

```
$ $HADOOP_HOME/bin/hadoop daemonlog -setlevel 127.0.0.1:50070 \
 org.apache.hadoop.hdfs.server.namenode.NameNode DEBUG
```

---

# Spark SQL Cheat Sheet

---

| SQL                                                      | Pyspark                                                               |
|----------------------------------------------------------|-----------------------------------------------------------------------|
| <code>select col1, ... from mutable</code>               | <code>dataFrame.select(col1, ...)</code>                              |
| <code>select count(*) from mytable</code>                | <code>dataFrame.count()</code>                                        |
| <code>select col1, col2, count(*) ... group by...</code> | <code>dataFrame.cube(col1, col2, ...)</code>                          |
| <code>select distinct ...</code>                         | <code>dataFrame.distinct()</code>                                     |
| <code>select ... where ...</code>                        | <code>dataFrame.filter(expression)</code>                             |
|                                                          | <code>dataFrame.groupBy(col1, ...)</code>                             |
| <code>select ... limit ...</code>                        | <code>dataFrame.limit(n)</code>                                       |
| <code>select ... order by ...</code>                     | <code>dataFrame.orderBy([col1, ...],<br/>ascending=True False)</code> |

***etc (MORE WORK ON THIS)***

*use quotes if columns are reserved words*

<https://spark.apache.org/docs/latest/api/python/pyspark.sql.html#pyspark.sql.DataFrame>



*Real-World Big Data in Action*

---

# Appendix

**BigData on Windows  
NOT CURRENTLY  
WORKING**

---

---

# Windows Prerequisites

---

## Install Cygwin (Windows 64-bit only)

- Provides a BASH shell to run scripts (not programs)
- Download from <https://cygwin.com/install.html>
- **Note that you must be running 64-bit Windows for Hadoop!**
- **DOESN'T WORK WITH HADOOP**

---

# Clone the Project Files

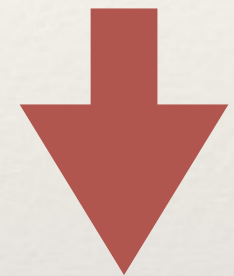
---

## Clone the Project files

```
$ cd $HOME/SPA_2016
```

```
$ git clone https://github.com/rozanski/bcs_spa16.git .
```

- don't miss out the dot at the end of the command!
- for Cygwin, add the flag `--config core.autocrlf=input` (avoids CRLF issues)
- create the remaining directories:





---

# Cygwin Setup

---

## Extra Packages When Installing Cygwin

- git, openssh

## Create a SPA 2016 User

- You need to add a user with a name without spaces (eg **spa16**)
- You can't do this from the UI since it demands a first and last name
- Run a Windows Command Prompt as Administrator

```
C:\Windows\System32>net user spa16 /add
```

- Log out to Windows, and log back in again as the spa16 user (you won't need to provide a password)
- Start a Cygwin Terminal
- Check you are running as the spa16 user

```
$ pwd
```

```
/home/spa16
```